

Notes			
Subject	Data Structure	Code	CIC-207
Community	SnapED Community		

Unit 3

Sorting

Sorting is a fundamental concept in computer science and everyday life. It refers to the process of arranging a collection of data items into a specific order, based on some chosen criterion. This order can be ascending (e.g., lowest to highest), descending (highest to lowest), alphabetical, chronological, or any other meaningful sequence.

Why is sorting important?

- Data analysis and retrieval: Sorted data is easier to analyze, visualize, and search. Imagine searching for a contact in an unsorted phone book vs. a sorted one!
- Efficiency: Sorting algorithms can significantly improve the performance of various tasks, from searching for information to merging datasets.
- Organization and clarity: Sorted data provides a clear structure and allows for easier identification of patterns and trends.

Types of sorting:

- Selection sort: iterates through the data and finds the minimum/maximum element, swapping it to the beginning/end. (Simple but slow)

- Bubble sort: repeatedly compares adjacent elements and swaps them if they're out of order. (Easy to understand but inefficient)
- Insertion sort: inserts each element into its rightful position in a sorted sub-sequence. (Efficient for small datasets)
- Merge sort: divides the data into halves, sorts each half recursively, and then merges them back in order. (Fast and efficient for large datasets)
- Quick sort: chooses a pivot element and partitions the data around it, recursively sorting each partition. (Fast but can be unstable)

Choosing the right sorting algorithm depends on factors like:

- Data size and type (numbers, strings, etc.)
- Desired order (ascending, descending, etc.)
- Performance requirements (speed, stability, etc.)
- Sorting can be done in-place or by creating a new copy of the data.
- Some algorithms are stable, meaning they preserve the original order of elements with equal values.
- The efficiency of a sorting algorithm is often measured by its time complexity, which indicates how the execution time grows with the data size.

Order

Order plays a crucial role in various data structures, impacting their efficiency and operations. Here's a breakdown of how order manifests in different structures:

Linear Data Structures:

- Arrays: Order is inherent, with elements accessed consecutively based on their index.
- Linked Lists: Order can be implicit (following pointers) or explicit (maintaining a "next" field).

- Queues: Order is strictly FIFO (First-In-First-Out), ensuring elements are processed in the order they are added.
- Stacks: Order is LIFO (Last-In-First-Out), with the most recently added element accessed first.

Tree Data Structures:

- Binary Search Trees (BSTs): Order is maintained through comparisons. Left subtree contains values less than the root, and right subtree holds values greater than the root. This enables efficient searching and sorting operations.
- Heaps: Order is based on a specific property (e.g., maximum or minimum value). Elements are arranged in a tree-like structure where parent nodes have higher/lower values than child nodes depending on the heap type.

Other Data Structures:

- Sets: Maintain unique elements but order is irrelevant. Operations like membership checking and union/intersection are efficient due to hashing.
- Graphs: Edges connect nodes, but their order doesn't matter. Traversal algorithms like DFS and BFS explore the graph based on the chosen starting point and edge connections.

Here's how order impacts operations:

- Searching: Ordered structures like BSTs enable efficient searching using comparisons. Unsorted structures require linear search, which is slower.
- Sorting: Sorting algorithms typically rely on the ability to compare and swap elements, which is facilitated by order.
- Efficiency: Maintaining order in certain structures can lead to faster operations like insertions, deletions, and updates compared to unordered structures.

Stability

The concept of stability in data structures goes beyond just maintaining order. It specifically refers to how a data structure preserves the relative order of elements with equal values during operations like insertion, deletion, or sorting.

Here's a deeper dive into stability:

What is a stable data structure?

A stable data structure ensures that if two elements have the same value and are initially in a specific order, that order will be preserved even after performing operations on the structure. Imagine two students with the same grade in a class roster. A stable data structure would keep them listed in the same order even after adding new students or sorting the roster.

Why is stability important?

- **Predictability:** Stable structures provide predictable behavior, simplifying code and debugging. You can be confident that elements with equal values will remain in their original relative positions.
- **Data integrity:** In certain situations, the relative order of equal elements might hold crucial information. Preserving this order through operations avoids data loss or misrepresentation.
- **Specific applications:** Some algorithms rely heavily on stability. For example, merging sorted lists or maintaining a consistent order of elements during iteration require stable structures.

Examples of stability in data structures:

- **Binary Search Trees (BSTs):** BSTs are generally considered stable because the order of elements with equal keys is preserved during insertion and deletion operations.
- **Doubly Linked Lists:** These lists maintain both forward and backward pointers, allowing them to remain stable during insertions and deletions anywhere in the list.

- Merge Sort: This sorting algorithm is inherently stable, as it merges pre-sorted sub-sequences while maintaining the relative order of elements with equal values.

Contrast with unstable structures:

- Insertion Sort: While efficient, insertion sort is not stable. Inserting an element with an existing value might change the relative order of other elements with the same value.
- Hash Tables: Hash tables store elements based on their hash values, which can lead to collisions and unpredictable order for elements with the same hash.

Choosing a stable or unstable structure depends on your specific needs and the operations you'll be performing. If maintaining the relative order of equal elements is crucial, then opting for a stable structure like a BST or a stable sorting algorithm like merge sort is recommended.

Selection Sort

There seems to be a slight misunderstanding here. While both selection sort and heap sort are sorting algorithms, they are distinct concepts and don't have variations like "straight" or "heap" within them.

Selection sort is a simple, comparison-based sorting algorithm. It repeatedly finds the minimum (or maximum) element in the unsorted portion of the data and swaps it with the first element. This process continues until the entire data is sorted. Selection sort is conceptually easy to understand but has a time complexity of $O(n^2)$, making it inefficient for large datasets.

Heap sort, on the other hand, utilizes a data structure called a heap to efficiently sort the data. A heap is a tree-based structure where each parent node has a higher (or lower) value than its children, depending on the type of heap (max-heap or min-heap). Heap sort builds the heap with the given data, then repeatedly extracts the minimum (or maximum) element from the root of the heap and rebuilds the heap

with the remaining elements. This process continues until all elements are extracted and sorted. Heap sort has a time complexity of $O(n \log n)$ and is significantly faster than selection sort for large datasets.

Insertion Sort and Shell Sort

Both insertion sort and shell sort are comparison-based sorting algorithms that work by iterating through the data and inserting each element into its rightful position in a sorted sub-array. However, they differ in their approach and efficiency.

Insertion Sort:

Concept: Imagine sorting cards in your hand. You compare each card with the ones before it and slide it into its rightful place if needed. Insertion sort does the same thing with data elements.

Steps:

1. Assume the first element is sorted.
2. Pick the next element (key).
3. Compare the key with the elements before it.
4. If the key is smaller, shift those elements one position to the right to make space.
5. Insert the key in the empty space.
6. Repeat steps 2-5 for all remaining elements.

Example:

Unsorted list: 8, 4, 6, 2, 1

1. Pass: Compare 4 with 8 (not smaller).
2. Pass: Compare 6 with 4 and 8 (not smaller).
3. Pass: Compare 2 with 4, 8, and 6 (smaller). Shift 6 and 8 to the right.

4. Pass: Insert 2 at its rightful position (before 6).
5. Pass: Continue comparing and shifting remaining elements.

Final sorted list: 1, 2, 4, 6, 8

Complexity: $O(n^2)$ in the worst case, meaning the time it takes to sort grows quadratically with the data size.

Straight Insertion Sort:

There is no "straight" insertion sort. It's simply the standard insertion sort as described above. Some variations exist, like gap insertion sort, which uses different intervals for comparison, but they are not commonly referred to as "straight insertion sort."

Shell Sort:

Concept: Shell sort is a hybrid of insertion sort and gap-based sorting. It starts with a larger gap between elements, allowing them to "jump" further and partially sort the data. Then, it gradually reduces the gap and applies insertion sort within smaller sub-sequences to refine the sorting.

Steps:

1. Choose a sequence of gaps (e.g., half, a quarter, eighth of the data size).
2. Sort elements within sub-sequences based on the chosen gap. This involves insertion sort-like comparisons and shifts.
3. Reduce the gap and repeat step 2 until the gap is 1.
4. Perform regular insertion sort on the entire data with a gap of 1.

Example:

Unsorted list: 8, 4, 6, 2, 1

1. Pass (gap 4): Compare and shift elements 4 positions apart (8 with 4, 6 with 2, etc.).
2. Pass (gap 2): Compare and shift elements 2 positions apart.
3. Pass (gap 1): Regular insertion sort to refine the sorting.

Final sorted list: 1, 2, 4, 6, 8

Complexity: $O(n \log n)$ on average, depending on the chosen gap sequence. This makes it generally faster than insertion sort for larger datasets.

In summary:

- Insertion sort is simple and easy to understand but slow for large datasets.
- Shell sort is a hybrid approach that improves efficiency by combining gap-based sorting with insertion sort.
- Choosing the best algorithm depends on the data size and specific requirements.

I hope this explanation clarifies the differences between insertion sort and shell sort, along with examples for better understanding. If you have any further questions about these algorithms or other sorting techniques, feel free to ask!

Exchange Sort (Bubble, quicksort)

You're right, both bubble sort and quicksort are considered "exchange sort" algorithms. This means they achieve sorting by repeatedly swapping elements within the data until the desired order is reached. However, they differ significantly in their approach and efficiency.

Bubble Sort:

- Concept: Imagine two children swapping marbles in their hands until they are in the right order. Bubble sort does the same for data elements, repeatedly swapping adjacent elements if they are in the wrong order.
- Steps:
 - Iterate through the data, comparing each pair of adjacent elements.
 - If the current element is larger than the next, swap them.
 - Repeat the process until no swaps occur in a full iteration, indicating the data is sorted.
-
- Example:
 - Unsorted list: 8, 4, 6, 2, 1
 - Pass 1: Compare and swap 8 and 4.
 - Pass 2: Compare and swap 6 and 4, 8 and 6.
 - Repeat passes until no swaps occur.
-
- Complexity: $O(n^2)$ in the worst case, making it inefficient for large datasets.
- Strengths: Simple to understand and implement.

Quicksort:

- Concept: Imagine dividing a deck of cards into smaller piles based on their values and then recursively sorting each pile. Quicksort uses a similar divide-and-conquer approach.
- Steps:
 - Choose a pivot element (often the first or last element).
 - Partition the data into two sub-sets: elements less than the pivot and elements greater than or equal to the pivot.
 - Recursively sort each sub-set using the same process.
 - Combine the sorted sub-sets back into a single sorted list.
-
- Example:

- Unsorted list: 8, 4, 6, 2, 1
- Choose pivot as 4.
- Partition data: [1, 2] and [8, 6].
- Recursively sort each sub-set.
- Combine sub-sets: [1, 2, 4, 6, 8].
-
- Complexity: $O(n \log n)$ on average, making it significantly faster than bubble sort for large datasets.
- Strengths: Efficient for large datasets, avoids unnecessary comparisons.
- Weaknesses: More complex to understand and implement than bubble sort, can be unstable in some implementations.

Merge Sort (External Sorting): Natural, Balanced, and Polyphase Merges

Merge sort is a versatile sorting algorithm that can handle both internal and external sorting. When dealing with massive datasets that cannot fit in memory, external sorting becomes necessary. This involves using temporary files on disk to perform the sorting process. Here's a breakdown of how merge sort adapts to external sorting and some specific variations you mentioned:

External Merge Sort:

1. **Data Chunking:** The large dataset is divided into smaller chunks that fit into memory. These chunks are called runs.
2. **Run Creation:** Each run is then sorted internally using a sorting algorithm like quicksort or merge sort itself. This creates multiple sorted files on disk.
3. **Merging Runs:** The sorted runs are then merged pairwise, combining them into larger sorted files. This process continues until a single sorted file remains.

Natural Merge:

- This is a simple form of external merge sort where runs are merged in adjacent pairs.
- Efficient for small datasets or when the number of available external storage devices is limited.
- However, for large datasets, it can lead to a lot of disk seeks and I/O operations, reducing efficiency.

Balanced Merge:

- Aims to minimize the number of disk seeks by merging runs that are farthest apart first.
- This reduces the average distance between elements being merged, leading to faster I/O performance.
- Requires keeping track of the positions of the runs and choosing the optimal pair for merging in each iteration.
- More complex to implement than natural merge but offers significant performance gains for large datasets.

Polyphase Merge:

- This variation utilizes multiple external storage devices (e.g., hard drives) to further improve efficiency.
- Instead of merging runs pairwise, it merges multiple runs simultaneously in phases.
- This can significantly reduce the number of merge passes compared to balanced merge and natural merge.
- Requires careful planning and synchronization to manage the merging process across multiple devices.

Choosing the Right Merge Variation:

- The choice depends on factors like:

- Data size
- Available external storage devices
- Performance requirements
- Complexity of implementation
-
- For small datasets, natural merge might be sufficient.
- For larger datasets, balanced merge offers better performance at the cost of increased complexity.
- Polyphase merge is ideal for massive datasets and scenarios with multiple external storage devices, but requires the most complex implementation.

Searching Algorithms

Searching algorithms are fundamental tools for efficiently finding specific elements within a collection of data. Here's a breakdown of the three you mentioned:

1. List Search:

This is the most general term encompassing any technique for finding an element in a list. It can involve various approaches like:

- Sequential Search: Traversing the list element by element until finding the target. Simple but slow for large lists.
- Binary Search: Only works with sorted lists. It repeatedly divides the search space in half based on comparisons, quickly narrowing down the target's location.
- Hashing: Uses a hash function to map elements to unique keys, allowing for constant-time lookup in a hash table (efficient for specific data types).
- Interpolation Search: Estimates the target's position based on its value and the surrounding elements in a sorted list, potentially reducing comparisons.

2. Sequential Search:

Also known as linear search, this algorithm iterates through each element in the list, comparing it to the target element. If a match is found, its position is returned. If not, the search continues until the end of the list.

Strengths:

- Easy to understand and implement.
- Works with any type of list, sorted or unsorted.

Weaknesses:

- Time Complexity: $O(n)$, meaning the search time grows linearly with the list size. This can be slow for large lists.
- Inefficient for sorted lists: Doesn't take advantage of the sorted order.

3. Binary Search:

This powerful algorithm only works with sorted lists. It works by repeatedly dividing the search space in half based on comparisons with the target element. If the target is smaller than the middle element, the search continues in the left half. If it's larger, the right half is considered. This process continues until the target is found or the search space shrinks to a single element.

Strengths:

- Time Complexity: $O(\log n)$, significantly faster than sequential search for large lists.
- Efficiently utilizes the sorted order of the list.

Weaknesses:

- Requires the list to be sorted.

- Can be slightly more complex to implement than sequential search.

Choosing the Right Search Algorithm:

The best choice depends on factors like:

- List size: For small lists, sequential search might be sufficient.
- List order: If the list is sorted, binary search is significantly faster.
- Data type: Hashing might be efficient for specific data types with efficient hash functions.
- Performance requirements: If speed is crucial, binary search or hashing might be preferred.

Hashing Methods and Collision Resolution: Finding the Right Key

Hashing is a powerful technique for storing and retrieving data efficiently. It works by mapping data elements to unique, fixed-size values called "hash codes" using a hash function. These hash codes then act as keys for accessing the data in a hash table, which is a specially designed data structure optimized for key-value lookups.

Here's a breakdown of hashing methods and collision resolution techniques:

Hashing Methods:

- Modulo Division: Divides the data element's value by a prime number and uses the remainder as the hash code. Simple but can lead to clustering for certain data sets.
- Bit Extraction: Selects specific bits from the data element's binary representation to create the hash code. Faster than modulo division but may not distribute elements evenly.
- Polynomial Hashing: Uses a polynomial function to map the data element to a hash code. More flexible than modulo division but computationally expensive.

Collision Resolution:

Collisions occur when different data elements map to the same hash code. This can happen due to limitations in the hash function or the size of the hash table. To handle collisions, several techniques can be used:

- Open Addressing:
 - Linear Probing: Checks for the next available slot in the hash table. Simple but can lead to clustering and performance degradation.
 - Quadratic Probing: Uses a quadratic function to probe for an empty slot. More complex but distributes elements more evenly.
 - Double Hashing: Uses a secondary hash function to probe for an empty slot. More complex but can be more efficient than linear or quadratic probing.
-
- Chaining: Stores colliding elements in a linked list attached to the colliding hash code slot. Easy to implement but can lead to long chains and performance problems.

Choosing the Right Hashing Method and Collision Resolution Technique:

The best choice depends on several factors:

- Data type and distribution: Different hashing methods work better for different data types and distributions.
- Performance requirements: Some methods are faster than others but may have higher collision rates.
- Memory limitations: Hashing utilizes additional memory for the hash table.
- Application requirements: Some applications require specific collision resolution properties, like guaranteed maximum search time.

UNIT-4

Disjoint Sets

Disjoint sets, also known as union-find data structures, represent a collection of sets that don't overlap or share elements. They are crucial for tasks like:

- Merging clusters: Grouping connected nodes in a graph.
- Detecting cycles: Finding loops in a graph or network.
- Path compression: Optimizing pathfinding algorithms.

There are two main ways to represent disjoint sets:

1. Linked List Representation:

- Each element in the set is represented by a node in a linked list.
- Each node has a pointer to the "parent" element, which represents the root of the set.
- Finding the set of an element involves traversing the list to the root.
- Union and find operations require updating parent pointers to keep the sets disjoint.

2. Forest Representation:

- Each set is represented by a tree, with the root as the set representative.
- Each node also has a rank (height of the subtree rooted at the node).
- Finding the set involves traversing the tree to the root.
- Union by rank merges two sets by attaching the tree with the smaller rank to the larger one, updating ranks as needed.

Advantages and Disadvantages:

- Linked List: Simple to implement, but find operations can be slow for deep sets.
- Forest Representation: Faster find operations, but union by rank can be more complex to implement.

Additional Techniques:

- Path compression: During find operations, update parent pointers to point directly to the root, improving future find performance.
- Union by size: Merges two sets by attaching the smaller set (by size) to the larger one, potentially improving performance compared to rank-based union.

Choosing the Right Representation:

The best representation depends on factors like:

- Frequency of find vs. union operations: If find operations are more frequent, a forest representation might be better.
- Data distribution: If sets tend to have similar sizes, union by size might be better.
- Complexity requirements: Linked list representation is simpler to implement, but forest representation with path compression can be much faster.

The union-find algorithm, also known as the disjoint-set data structure, is a powerful tool for managing a collection of non-overlapping sets. It efficiently performs two key operations:

- Find: Determines which set a specific element belongs to.
- Union: Merges two sets into a single set.

This algorithm is used in various applications, including:

- Clustering: Grouping connected components in a graph.

- Network analysis: Detecting cycles or connected components in a network.
- Pathfinding: Optimizing pathfinding algorithms by identifying connected regions.
- Image segmentation: Grouping pixels with similar properties into distinct regions.

Union Find Algorithm

There are two main approaches to implementing the union-find algorithm:

1. Union by Rank:

- Each set is represented as a tree, with the root representing the set.
- Each node has a rank, which is the height of the subtree rooted at that node.
- Find: Traverse the tree to the root, updating parent pointers for path compression.
- Union: Attach the tree with the smaller rank to the one with the larger rank. Update ranks as needed to maintain tree heights.

This approach balances the heights of trees in the union operation, leading to faster find operations on average.

2. Union by Size:

- Similar to union by rank, but sizes of sets are used instead of ranks.
- Find: Traverse the tree to the root, updating parent pointers for path compression.
- Union: Attach the smaller set (by size) to the larger one. Update sizes as needed.

This approach is generally simpler to implement than union by rank, but it can lead to unbalanced trees and slower find operations for certain data distributions.

Additional Techniques:

- Path Compression: During find operations, directly update parent pointers to point to the root of the set, improving future find performance.
- Splitting: Occasionally split large trees to maintain a desired maximum height and prevent slowdowns.

Choosing the Right Implementation:

- Union by rank generally offers better performance for most cases, especially with path compression.
- Union by size can be easier to implement and might be preferable for specific data distributions where rank-based union performs poorly.

Graphs and Graph Representation

Graphs are powerful tools for representing relationships between objects. They consist of two main components:

- Vertices: Represent individual entities (e.g., people, cities, keywords).
- Edges: Represent connections or relationships between vertices (e.g., friendship, transportation routes, semantic links).

Graph representation refers to the techniques used to store and manipulate graphs in computer memory. Choosing the appropriate representation is crucial for efficient algorithms and data analysis. Here are some common options:

1. Adjacency Matrix:

- A two-dimensional matrix where rows and columns represent vertices.
- Each cell indicates the presence or absence of an edge between the corresponding vertices.

- Efficient for checking if two vertices are connected but can be memory-intensive for sparse graphs (many missing edges).

2. Adjacency List:

- Each vertex has a list of its neighboring vertices.
- More memory-efficient than adjacency matrices for sparse graphs.
- Requires additional processing to determine if two arbitrary vertices are connected.

3. Edge List:

- Explicitly stores each edge as a pair of vertices.
- Simple to implement but inefficient for searching for neighbors or performing graph operations.

4. Incidence Matrix:

- Similar to an adjacency matrix, but rows represent vertices and columns represent edges.
- Each cell indicates if a vertex is incident to an edge (connected by the corresponding edge).
- Useful for analyzing the properties of edges and their relationships to vertices.

5. Object-Oriented Representation:

- Each vertex and edge is represented as a separate object with associated properties.
- Offers flexibility and encapsulation but can be more complex to implement.

Choosing the Right Representation:

The best representation depends on factors like:

- Graph density: Sparse graphs benefit from adjacency lists, while dense graphs might be better suited for adjacency matrices.
- Operations frequently performed: If checking for connections is crucial, an adjacency matrix might be efficient. If iterating through neighbors is common, an adjacency list might be better.
- Memory constraints: Adjacency lists and edge lists are generally more memory-efficient than adjacency matrices.
- Directed vs. undirected graphs: Representation needs to adapt to the directionality of edges, if applicable.
- Weighted graphs: Edges can have associated values (weights), requiring additional storage and handling.

Graph Traversals

Graph traversals are fundamental algorithms used to systematically visit all vertices or edges in a graph. They are essential for various tasks, including:

- Finding shortest paths: BFS is efficient for finding the shortest path between two vertices.
- Detecting cycles: DFS can effectively identify cycles (closed loops) within the graph.
- Topological sorting: Ordering vertices in a directed acyclic graph based on their dependencies.
- Connected components: Identifying groups of vertices connected to each other.

Two popular graph traversal algorithms are:

1. Breadth-First Search (BFS):

- Concept: Imagine exploring a maze by visiting all neighboring nodes level by level before moving deeper.
- Implementation:
 - Use a queue to store unvisited vertices.
 - Dequeue a vertex, visit it, and add its unvisited neighbors to the queue.
 - Repeat until the queue is empty.
-
- Properties:
 - Finds shortest paths between vertices.
 - Efficient for dense graphs.
 - Can be used for level-order traversal of trees.
-

2. Depth-First Search (DFS):

- Concept: Imagine exploring a maze by following a single path as far as possible before backtracking.
- Implementation:
 - Use a stack to track visited vertices.
 - Push a vertex onto the stack and explore its unvisited neighbors recursively.
 - Backtrack if no unvisited neighbors are found.
 - Repeat until the stack is empty.
-
- Properties:
 - Efficient for finding paths in sparse graphs.
 - Can be used for detecting cycles and topological sorting.
 - May visit the same vertex multiple times in some cases.
-

Implementation Considerations:

- Both algorithms can be implemented iteratively or recursively.
- Marking visited vertices is crucial to avoid infinite loops.
- Different data structures (queues, stacks) impact efficiency and memory usage.

Choosing the Right Algorithm:

The best choice depends on the specific needs:

- Shortest paths: Use BFS.
- Cycles or topological sorting: Use DFS.
- Dense vs. sparse graphs: Consider efficiency trade-offs.
- Memory limitations: Choose an iterative implementation if necessary.

Further Exploration:

- There are variations of BFS and DFS (e.g., bidirectional BFS) with specific advantages.
- Other graph traversal algorithms exist for different purposes (e.g., Dijkstra's algorithm for finding shortest weighted paths).

Both Minimum Spanning Tree (MST) and Shortest Path algorithms deal with finding optimal paths in graphs, but they address different objectives and have distinct approaches:

Minimum Spanning Tree (MST):

- Objective: Find a subset of edges in a connected, weighted graph that connects all vertices with the minimum total weight, while ensuring no cycles.
- Applications: Network design, clustering, data compression, etc.
- Popular algorithms: Prim's algorithm, Kruskal's algorithm, Boruvka's algorithm.

- Key point: Minimizes the total weight of the spanning tree, not the individual edge weights between specific vertices.

Shortest Path Algorithms:

- Objective: Find the shortest path (minimum weight sequence of edges) between a source and a destination vertex in a weighted graph.
- Applications: Navigation, routing, scheduling, etc.
- Popular algorithms: Dijkstra's algorithm, Bellman-Ford algorithm, A* search.
- Key point: Focuses on the shortest path between two specific vertices, not necessarily minimizing the overall weight of the connected network.

Key Differences:

- Scope: MST considers all vertices and minimizes the total weight of the tree, while shortest path focuses on a specific pair of vertices and minimizes the weight of the path between them.
- Cycles: MST algorithms ensure no cycles, while shortest path algorithms may sometimes involve cycles (e.g., when the destination is reachable through a shorter path with a loop).
- Applications: MSTs are useful for resource allocation and network design, while shortest paths are crucial for navigation, planning, and optimization tasks.

Choosing the Right Algorithm:

- If you need to connect all vertices in a network with minimal total cost and avoid cycles, use an MST algorithm.
- If you need to find the shortest path between two specific vertices, regardless of cycles, use a shortest path algorithm.

